


Q00 native LoRA: output projection vs attention updates

Primary: fixed LR 3e-5



Large dots: mean of 3 corpus/optimizer repetitions; whiskers: ± 1 repetition SD, not CIs.

Small dots: individual repetitions. Labels include trainable parameters; capacities are unequal.

512 shared evaluation episodes, previously examined Q00. F0/BOTH are reused references. Both procedures select LR 3e-5.