




Does training in the evaluated space improve forecasting?



Development diagnostic: two previously evaluated sources, one optimizer seed (12000), LR 3e-5, 200 updates.

Initial gradient norms are matched; AdamW updates and subsequent clipping are not matched. This is not a novel PEFT method.

Primary: SORT NORM_ALIGNED minus RAW_ALIGNED, divided by SORT F0. Date-block CIs exclude training uncertainty.